

15 seconds (cocktail party / casual): ArtHound turns the messy production data scattered across studios' tools into one clean, usable layer — so producers can finally answer the questions they've been answering with spreadsheets and gut feel. We ship the questions we think matter most, but the real value is that your data is finally yours to use. 30 seconds (warm intro / first investor touch): Game studios run production on a patchwork of Airtable, ShotGrid, vendor portals, and spreadsheets. The data is fragmented, inconsistent, and unusable for any question that crosses tools — which is most of the questions that matter. ArtHound is a canonical production data layer that normalizes this mess into clean structured data with stable IDs and per-studio mapping. We ship opinionated capabilities on top — work generation, scenario planning, agent-driven queries — but the foundational value is that studios can finally use their own data, in our tools or theirs. I built the prototype of this for VALORANT at Riot. Now I'm productizing it. 60 seconds (full pitch / pitch meeting): Game production at scale runs on fragmented data. A typical studio's asset information lives across Airtable, ShotGrid, vendor systems, time tracking, and finance — none of which talk to each other in a structured way. This means the questions leadership actually cares about — what's our real burn rate, what happens if we shift work to Singapore, where are we late and why — get answered with manual spreadsheet pulls or not at all. It's also why studios can't deploy AI agents against their production data, even though every studio is trying. ArtHound is the canonical data layer underneath all of this. We sync from existing source systems, normalize into a clean schema with stable global IDs, and maintain parity so studios don't have to switch tools. On top of that layer, we ship capabilities studios need immediately — work generation, scenario planning, agent-driven queries against their data. But the durable value is the layer itself: studios can use it through our features, through their own BI tools, or through whatever agent framework they choose. I led production for VALORANT cosmetics at Riot, where I built the internal version of this for a billion-dollar live-service operation across three regions and multiple outsource partners. The architecture I'm productizing now is what made that operation possible. The market timing is that every studio is trying to deploy AI on production data and discovering their data isn't ready — we make it ready, and we sell the capabilities that prove the value.

A few notes on the pitches: The 60-second version assumes you're talking to someone who'll engage with the category-creation framing. Most pre-seed investors will. Some will not, and for those the 30-second version with a sharper feature anchor (work generation + scenario planning) lands better. Worth sharpening before you actually use any of these: your founder credential line. "Senior Manager who ran VALORANT Expressions and built the production data infrastructure for it" is more specific than what I wrote and lands harder. The numbers — three regions, ~\$100M esports revenue context, etc — are the parts that make investors lean in, but they need to be deployed precisely and not as a list. One concrete number embedded in the narrative is more powerful than three numbers stacked together. The pitch doesn't yet handle the "why now" question well. Right now it's implicit (AI agents need clean data). Worth deciding whether to make that explicit in the elevator pitch or save it for the follow-up conversation. Explicit is more memorable; implicit is less reductive. Your call based on the audience. The pitch also doesn't address the obvious investor objection: "isn't this what ShotGrid/ftrack should be doing." The answer is that they can't redesign their data model without breaking customers, but you don't have time to make that argument in 60 seconds. Have it loaded for the question, which will come within 5 minutes of you finishing the pitch.